



# Component-based parallelization of multi-stencil programs

Hélène Coullon, Christian Perez

**RESEARCH  
REPORT**

**N° 7003**

April 2015

Project-Teams Avalon





## Component-based parallelization of multi-stencil programs

Hélène Coullon<sup>\*†</sup>, Christian Perez<sup>‡</sup>

Project-Teams Avalon

Research Report n° 7003 — April 2015 — 9 pages

**Abstract:** abstract

**Key-words:** No keywords

---

\* Footnote for first author

† Shared foot note

‡ Footnote for second author

**RESEARCH CENTRE  
GRENOBLE – RHÔNE-ALPES**

Inovallée  
655 avenue de l'Europe Montbonnot  
38334 Saint Ismier Cedex

**Exemple de document  
utilisant le style  
rapport de recherche  
Inria**

**Résumé :** Resume en francais

**Mots-clés :** mots-cles

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Computational model of multi-stencil programs</b>	<b>3</b>
2.1	Definitions . . . . .	3
2.2	Parallelization . . . . .	5
<b>3</b>	<b>Component model and transformation</b>	<b>9</b>
3.1	A primitive component model . . . . .	9
3.2	Control components . . . . .	9
3.3	Transformation to a parallel assembly . . . . .	9
<b>4</b>	<b>Component Stencil Language</b>	<b>9</b>
4.1	CSL language and example . . . . .	9
4.2	CCSL compiler and example . . . . .	9
4.3	Evaluation . . . . .	9
<b>5</b>	<b>Related work</b>	<b>9</b>
<b>6</b>	<b>Conclusion</b>	<b>9</b>

## 1 Introduction

## 2 Computational model of multi-stencil programs

To numerically solve a set of PDEs, iterative methods are frequently used to approximate the solution by a step by step phenomena. Thus, the continuous time and space domains are discretized so that a set of numerical computations are iteratively (time discretization) applied on a mesh (space discretization). In other words, the PDEs are transformed to a set of numerical computations applied at each time step on all elements of the discretized space domain. Among the numerical computations is found a set of numerical schemes, also called *stencil computations*. A formal definition of a *stencil program*, and the parallelization of such program, are presented in this section.

### 2.1 Definitions

A mesh  $\mathcal{M}$  defines the discretization of the continuous space domain  $\Omega$  of a set of PDEs and is defined as followed.

**Definition** A mesh is a connected undirected graph  $\mathcal{M} = (V, E)$ , where  $V$  is the set of vertices and  $E$  the set of edges. The set of edges  $E$  of a mesh  $\mathcal{M} = (V, E)$  does not contain bridges.

**Definition**  $D_i$  is a set of elements of a mesh  $\mathcal{M} = (V, E)$ , constructed by a function  $domain_i$  which defines a precise association between  $V$  and  $E$ ,  $domain_i : V \times E \rightarrow D_i$ .

For example, the set of cells  $D_0$  in a Cartesian 2D mesh could be defined by exactly four vertices and four edges connected as a cycle. But we could also define another set of elements  $D_1$  as the simple set of vertices  $V$  etc.

A mesh can be structured (as Cartesian or curvilinear meshes), unstructured, regular or irregular (without the same topology for each element) and hybrid.

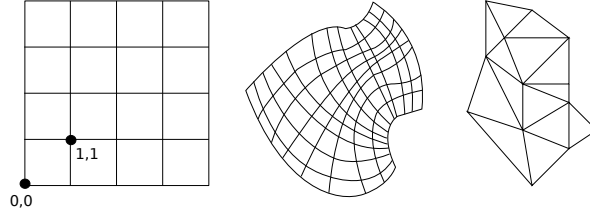


Figure 1: From left to right, Cartesian, curvilinear and unstructured meshes.

**Definition** The discretization of the continuous time domain  $\mathcal{T}$  is denoted  $T$  such that  $\forall t_i, t_{i+1} \in T, \exists \Delta t \in \mathbb{R}, t_{i+1} = t_i + \Delta t$ . Thus,  $T$  is responsible for the iteration time steps of the numerical simulation.

In a numerical simulation a set of data, or quantities, are applied onto the mesh and represent the set of values to compute, or to use, for computation.

**Definition** The set of data applied on the mesh is denoted by  $\Delta$ , such that  $\delta \in \Delta$  is a function which associates each element  $d \in D_i$  to a value  $v \in V$ ,  $\delta : D_i \rightarrow V$ .

One can notice that in applied mathematics, the signature of  $\delta$  would be  $\delta : E_i \times T \rightarrow V$ , however when programming a numerical simulation it is not wise to store all values of each time iteration.

**Definition** A numerical expression  $\exp$  is a function which represents how to compute, for an element  $d \in D_i$ , a data  $w \in \Delta$  (written data) with a set  $R \subset \Delta$  of input data (read data),  $\exp : R \times D_i \rightarrow w \times D_i$ .

**Definition** A computation  $c$  of a numerical simulation is defined as  $c(R, w, D, \exp)$ , where  $R \subset \Delta, w \in \Delta$  and  $\exp$  a numerical expression  $\exp : R \rightarrow w$ .  $D$  is one of the subsets  $D_i \subset \mathcal{M}$ , such that  $w : D \rightarrow V$ .

It has to be noticed that at each time iteration, all the elements of a mesh are computed. However, it happens that the computation of the mesh elements is splitted in different computations (for example the computation of the physical border). In this case additional  $D_i$  can be specified for the mesh  $\mathcal{M}$ .

**Definition** The set of  $n$  ordered computations of a numerical simulation is denoted  $\Gamma = [c_i]_{0 \leq i \leq n-1}$ , such that  $\forall c_i, c_j$  with  $i \leq j$ ,  $c_i$  is computed before  $c_j$ , and  $c_j$  can be computed only when  $c_i$  is finished.

**Definition** Finally, a *multi-stencil program* is defined by the quadruplet  $\mathcal{MSP}(T, \mathcal{M}, \Delta, \Gamma)$ .

**Definition** The neighborhood  $\mathcal{N}$  of an element  $d \in D_i$  is a function to define a set of elements in any  $D_k \subset \mathcal{M}$ ,  $\mathcal{N} : D_i \rightarrow D_k \times D_k \times \dots$ .

The function  $\mathcal{N}$  is also sometimes called the *stencil shape*, or the *stencil* in applied mathematics.

A computation  $c \in \Gamma$  can be of two different types. The first type is called a *stencil computation*.

**Definition** A *stencil computation* is defined as a quintuplet  $s(R, w, D, \exp, \mathcal{N})$ , where  $R \subset \Delta, w \in \Delta$ ,  $D$  is one the subsets  $D_i$  and  $w : D \rightarrow V$ .

In a stencil computation  $s$ ,  $\forall d \in D$ , the stencil numerical expression  $\exp$  is applied such that  $w(d) = \exp(R(d), R(\mathcal{N}(d)))$ . In this work, a stencil computation  $s(R, w, D, \exp, \mathcal{N})$  always verifies  $R \cap w = \emptyset$ , otherwise an implicit numerical scheme has to be solve which is over the scope of this paper.

Figure 2 gives an example of a stencil computation  $s(R, w, D, \exp, \mathcal{N})$ , where  $\mathcal{M}(V, E)$  is a two dimensional Cartesian mesh. A single domain  $D$  is defined in this example and is composed of cells formed by a cycle of four vertices  $v \in V$  and four edges  $e \in E$ . Furthermore, in this example  $R = \{A\}$ ,  $w = B$ , and for  $(x, y) \in D$  the neighborhood function is

$$\mathcal{N} : (x, y) \rightarrow \{(x, y + 1), (x, y - 1), (x + 1, y), (x - 1, y)\}.$$

Finally, the numerical expression of this example is

$$\exp(A(x, y), A(\mathcal{N}(x, y))) = B(x, y) = A(x, y) + (A(x, y + 1) + A(x, y - 1) + A(x + 1, y) + A(x - 1, y)) / 4.$$




Figure 2: Example of a stencil computation.

Finally, the second type of numerical computation is a local computation.

**Definition** A local computation is a quadruplet  $l(R, w, D, \exp)$ , where  $e$  does not involve a neighborhood function  $\mathcal{N}$ .

## 2.2 Parallelization

Multi-stencil mesh-based numerical simulation can be parallelized in various ways and is an interesting kind of application to take advantage of modern heterogeneous HPC architectures, mixing clusters, multi-cores CPUs, vectorization units, GPGPU and many-core accelerators.

**Coarse-grain data parallelism.** In a data parallelization technique, the idea is to split the data on which the program is computed in balanced sub-parts, one for each available resource. The same sequential program can afterwards be applied on each sub-part simultaneously, with some additional synchronizations between resources to update the data not computed locally, and thus to guarantee a correct result.

More formally, the data parallelization of a multi-stencil program  $\mathcal{MSP}(T, \mathcal{M}, \Delta, \Gamma)$  consists in, first, a partitioning of the mesh  $\mathcal{M}$  in  $p$  balanced sub-meshes (for  $p$  resources)  $\{\mathcal{M}_0, \dots, \mathcal{M}_{n-1}\}$ . This step can be performed by an external graph partitionner  $\parallel$  and is not addressed by this paper. As a data is mapped onto the mesh, the set of data  $\Delta$  is partitionned the same way than the mesh in  $\{\Delta_0, \dots, \Delta_{n-1}\}$ . The second step of the parallelization is to identify in  $\Gamma$  the needed synchronizations between resources to update data, and thus to build a new ordered list of computations  $\Gamma_{data}$ .

**Definition** For  $n$  the number of computations in  $\Gamma$ , and  $i < j < n$ , a *synchronisation* is needed between  $c_i$  and  $c_j$ , denoted  $c_i \ll c_j$ , if  $c_j = s_j(R_j, w_j, D_j, \text{exp}_j)$ , and  $w_i \subset R_j$ . Moreover, the data to update is  $w_i \cap R_j = w_i$ .

Actually, a synchronization can only be needed by the data read in a stencil computation, and only if this data has been modified before, which means that it has been written before. This synchronization is needed because the neighborhood function  $\mathcal{N}$  of the stencil computation involves values computed on different resources.

**Definition** A synchronization between two computations  $c_i \ll c_j$  is defined as a specific computation  $\text{update}(w_i \cap R_j) = \text{update}(w_i)$ .

**Definition** The concatenation of two ordered lists of respectively  $n$  and  $m$  computations  $l_1 = [c_i]_{0 \leq i \leq n-1}$  and  $l_2 = [w_i]_{0 \leq i \leq m-1}$  is denoted  $l_1 \cdot l_2$  and is equal to a new ordered list  $l_3 = [c_0, \dots, c_{n-1}, w_0, \dots, w_{m-1}]$ .

**Definition** From the ordered list of computation  $\Gamma$ , a new ordered list  $\Gamma_{data}$  is obtained from the call  $\Gamma_{data} = T_{data}(\Gamma, 0)$ , where  $T_{data}$  is the recursive function defined as

$$T_{data}(\Gamma, i) = \begin{cases} [\Gamma[i]] \cdot T_{data}(\Gamma, i+1) & \text{if } \forall j \leq i, c_j \not\ll c_i \\ [\text{update}(\{w_k\}), \Gamma[i]] \cdot T_{data}(\Gamma, i+1) & \forall k \leq i, c_k \ll c_i \\ [] & \text{if } i = |\Gamma|. \end{cases}$$

The construction of  $\Gamma_{data}$  implies that if a synchronization is needed it is always proceeded before the computation which needs it. However, an optimization is possible to combine synchronizations together, and thus to synchronize more than one data at a time. If the number of synchronizations is reduced, final performance of the program will be better. For this reason, we define a reduction to apply to  $\Gamma_{data}$ , and which is performed as a post-transformation.

**Proposition 2.1** Denoting  $c_k$  and  $c_j$ , with  $k < j$ , two computations of the ordered list  $\Gamma_{data}$  for which a synchronization is needed. If  $c_i$ , with  $i < j$  is the computation such that  $c_i \ll c_j$  and such that  $\text{update}(w_i)$  is performed before  $c_j$  in  $\Gamma_{data}$ . And if the needed synchronization before  $k$  is  $\text{update}(w_l)$ , where  $l < k$ . Then, if  $i < k$ , a single synchronization phase is needed before  $c_k$  and is  $\text{update}(w_l \cup w_i)$ .

**Proof** By definition, the synchronization of  $w_i$  can be performed between  $c_i$  and  $c_j$ . As  $i < k$  the computation  $c_i$  is not performed between  $k$  and  $j$ . As a result, the synchronization of  $w_i$  can be performed between  $c_i$  and  $c_k$ .

In other words, for each synchronization in  $\Gamma_{data}$ , if the computation  $c_i$  from which  $c_j$  is synchronized occurs before the last synchronization, the two synchronizations can be reduced to one with the union of the data to update.

The final step of this parallelization is to run  $\Gamma_{data}$  on each resource. Thus, for each resource  $0 \leq k \leq p-1$  a multi-stencil program defined by

$$\mathcal{MSP}_k(T, \mathcal{M}_k, \Delta_k, \Gamma_{data}), \quad (1)$$

is runned.

We denote this parallelization technique a coarse-grain data parallelization in contrast with the same technique applied to the finer level of a single computation. In this case, for a computation  $c(R, w, D, \text{exp})$ , the local domain  $D$  is partitionned for  $p$  resources  $\{D_0, \dots, D_{p-1}\}$ , and each resource  $k$  is responsible for a sub-computation  $c_k(R, w, D_k, \text{exp}')$ . This finer data parallelization technique is not directly addressed by the work presented in this paper, but is exhibited, as it will be explained later.



**Task parallelism.** A task parallelization technique is a technique to transform a program as a dependency graph of different tasks. A dependency graph exhibits parallel tasks, or on the contrary sequential execution of tasks. In a multi-stencil program a task is defined as a computation  $c(R, w, D, \text{exp})$ . As a result, the dependency graph of a multi-stencil program represents the dependencies between local and stencil computations of  $\Gamma$ .

**Definition** For two computations  $c_i(R_i, w_i, D_i, e_i)$  and  $c_j(R_j, w_j, D_j, e_j)$ , with  $i \leq j$ , it is said that  $c_j$  is dependant from  $c_i$ , denoted  $c_i \prec c_j$ , if  $w_i \cap R_j \neq \emptyset$ . In this case,  $c_i$  has to be computed before  $c_j$ . The binary relation  $c_i \prec c_j$  represents a *dependency*.

**Proposition 2.2** *The binary relation  $\prec$  is not transitive.*

**Proof** Considering three computations  $c_0(R_0, w_0, D_0, e_0)$ ,  $c_1(R_1, w_1, D_1, e_1)$  and  $c_2(R_2, w_2, D_2, e_2)$ , where  $w_0 \cap R_1 \neq \emptyset$  and  $w_1 \cap R_2 \neq \emptyset$ . In this case two dependencies can be extracted  $c_0 \prec c_1$  and  $c_1 \prec c_2$ . However, the dependency  $c_0 \prec c_2$  is not true as  $w_0 \cap R_2 = \emptyset$ .

It has been proved that the binary relation  $\prec$  is not transitive. This property is due to the fact that  $\prec$  is defined as a data dependency between computations. A dependency chain, however, creates another type of dependency only due to time.

**Definition** For two computations  $c_i(R_i, w_i, D_i, e_i)$  and  $c_j(R_j, w_j, D_j, e_j)$ , a time dependency is denoted  $c_i \blacktriangleleft c_j$  and means that  $c_i$  has to be computed before  $c_j$ .

**Proposition 2.3** *For two computations  $c_i(R_i, w_i, D_i, e_i)$  and  $c_j(R_j, w_j, D_j, e_j)$  such that  $c_i \prec c_j$ ,  $c_i \blacktriangleleft c_j$  is verified.*

**Proof** By definition  $c_i \prec c_j$  means that  $w_i \cap R_j \neq \emptyset$  and that  $c_i$  has to be computed before  $c_j$ .

**Proposition 2.4** *The binary relation  $\blacktriangleleft$  is transitive.*

**Proof** Considering three computations  $c_0(R_0, w_0, D_0, e_0)$ ,  $c_1(R_1, w_1, D_1, e_1)$  and  $c_2(R_2, w_2, D_2, e_2)$ , where  $c_0 \blacktriangleleft c_1$  and  $c_1 \blacktriangleleft c_2$ .  $c_0$  is computed before  $c_1$  and  $c_1$  is computed before  $c_2$ , as a result  $c_0$  is computed before  $c_2$  and the relation  $c_0 \blacktriangleleft c_2$  is verified.

**Definition** A directed acyclic graph (DAG)  $G(V, A)$  is a graph where the edges are directed from a source to a destination vertex, and where following the direction of edges, no cycle can be found from a vertex  $u$  to itself. A directed edge is called an arc, and for two vertices  $v, u \in V$  an arc from  $u$  to  $v$  is denoted  $(u, v) \in A$ .

From an ordered list of computations, a directed dependency graph  $\Gamma_{task}(V, A)$  can be built finding all pairs of computations  $c_i(R_i, w_i, D_i, e_i)$  and  $c_j(R_j, w_j, D_j, e_j)$ , with  $i < j$ , such that  $c_i \blacktriangleleft c_j$ . Those pairs of relations can be found from pairs of data dependencies  $c_i \prec c_j$ .

**Definition** From an ordered list of computation  $c(R, w, D, \text{exp})$ , a directed dependency graph  $\Gamma_{task}(V, A)$  is obtained from the call  $T_{task}(\Gamma, 0)$ , where  $T_{task}$  is the recursive function

$$T_{task}(\Gamma, i, \Gamma_{task}(V, A)) = \begin{cases} (V, A) & \text{if } i = |\Gamma| \\ T_{task}(\Gamma, i + 1, \Gamma_{task}(V \cup c_i, A \cup (c_k, c_i))), \forall k \leq i, c_k \prec c_i & \text{if } i < |\Gamma| \end{cases}$$

This constructive function is possible because the input is an ordered list. Actually, if  $c_k \prec c_i$  then  $k < i$ . As a result,  $c_k$  is already in  $V$  when the arc  $(c_k, c_i)$  is built.

**Proposition 2.5** *The directed graph  $\Gamma_{task}$  is an acyclic graph.*

**Proof**  $\Gamma_{task}$  is built from  $\Gamma$  which is an ordered and sequential list of computations. Moreover, each computation of the list  $\Gamma$  is associated to a vertex of  $V$ , even if the same computation is represented more than once in  $\Gamma$ . As a result it is not possible to go back to a previous computation and to create a cycle.

Using the function  $T_{task}$  to build  $\Gamma_{task}$ , however, duplication of dependencies may occur because of the transitivity of the relation  $\blacktriangleleft$ . Actually, as the relation  $\prec$  verifies the relation  $\blacktriangleleft$ , and as  $\blacktriangleleft$  is transitive, if for three computations  $c_k(R_k, w_k, D_k, e_k)$ ,  $c_i(R_i, w_i, D_i, e_i)$  and  $c_j(R_j, w_j, D_j, e_j)$ , with  $k < i < j$  and where  $c_k \prec c_i$ ,  $c_i \prec c_j$  and  $c_k \prec c_j$ , then the *pseudo-transitivity*  $c_0 \blacktriangleleft c_2$  is verified because of the relations  $c_0 \blacktriangleleft c_1$  and  $c_1 \blacktriangleleft c_2$ . However,  $c_0 \blacktriangleleft c_2$  is also directly represented by the binary relation  $c_0 \prec c_2$ . As a result, a duplication of the dependency is created in  $\Gamma_{task}$ . In Figure 3 an example of duplication is given for the relation  $c_1 \blacktriangleleft c_4$ .




Figure 3: Useless duplication of the dependency  $c_1 \blacktriangleleft c_4$

Another view of the relations  $\prec$  and  $\blacktriangleleft$  is that  $c_i \blacktriangleleft c_j$  in  $\Gamma_{task}$  is a path from  $c_i$  to  $c_j$  of any size, while  $c_i \prec c_j$  is a path of size 1 from  $c_i$  to  $c_j$ . And in any case a path in the graph represents a time dependency between computations. Scheduling a graph dependency which do not contains duplication of information is easier. For this reason, we define a reduction to apply to  $\Gamma_{task}$ , and which is performed as a post-transformation. In the example of Figure 3 the dependency represented by the path  $c_1 \blacktriangleleft c_2 \blacktriangleleft c_3 \blacktriangleleft c_4$  gives more information than the one directly represented by  $c_1 \prec c_4$ . The proposed reduction is based on this property.

**Definition** The dependency DAG  $\Gamma_{task}(V, A)$  is reduced to the dependency graph  $\Gamma_{task}(V, A')$ , where  $A' \subseteq A$ . A direct arc  $(c_i, c_j)$ ,  $i < j$ , of  $A$ , which means that  $c_i \prec c_j$ , is included in  $A'$  if no longer path  $c_i \blacktriangleleft c_j$  exists in  $A$  because of transitivity.

**Hybrid parallelism.** It is also possible to combine coarse-grain data and task parallelization techniques to get hybrid parallelism, sometimes more efficient on hybrid architectures. A coarse-grain data parallelism creates for  $k$  resources  $k$  multi-stencil computations  $\mathcal{MSP}_k(T, \mathcal{M}_k, \Delta_k, \Gamma_{data})$ , where  $\Gamma_{data}$  is built from an ordered list of computations and is itself an ordered list with the addition of needed synchronizations. On the other hand, the task parallelization technique builds, from an ordered list of computation, a graph dependency  $\Gamma_{task}$ . As a result  $T_{data}$  and  $T_{task}$  can be composed to build  $\Gamma_{hybrid}$  such as  $k$  multi-stencil programs

$$\mathcal{MSP}_k(T, \mathcal{M}_k, \Delta_k, \Gamma_{hybrid})$$

are responsible for an hybrid parallelization of a multi-stencil program. The set of computation  $\Gamma_{hybrid}$  is a DAG of dependencies and synchronizations and can be built from the call to

$$T_{task}(T_{data}(\Gamma, 0), 0, \Gamma_{hybrid}).$$

## 3 Component model and transformation

### 3.1 A primitive component model

Definitions of a primitive component (ports, interfaces), a component instantiation, an assembly + simple semantics on the behavior of a primitive component. Cite L2C as an implementation of such a model + introduction of notations (use,provide,use-multiple,sync).

### 3.2 Control components

Definition of three control components, and of the kernel component. Explain the behavior.

### 3.3 Transformation to a parallel assembly

Transformation from a parallel algorithm using the directives: PARALLEL, PARALLEL SECTIONS, SECTION and FORALL, to a component assembly

## 4 Component Stencil Language

### 4.1 CSL language and example

### 4.2 CCSL compiler and example

### 4.3 Evaluation

## 5 Related work

stencil compilation: Pochoir, PATUS etc.

stencil programs: Lizst, OP2

control in components: X-MAN (The New Component Model), STCM, Kell-calculus

## 6 Conclusion



**RESEARCH CENTRE  
GRENOBLE – RHÔNE-ALPES**

Inovallée  
655 avenue de l'Europe Montbonnot  
38334 Saint Ismier Cedex

Publisher  
Inria  
Domaine de Voluceau - Rocquencourt  
BP 105 - 78153 Le Chesnay Cedex  
[inria.fr](http://inria.fr)

ISSN 0249-6399